



VIDEO

Q&A WITH SPEAKER [SPEAKER BIO](#)

Ketan Umare
Software engineer
Lyft

Ketan Umare is a senior staff software engineer at Lyft a the Flyte project. Before Flyte he worked on ETA, routing infrastructure at Lyft. He is also the founder of Flink Kub-operator and contributor to Spark on K8s. Prior to Lyft he founding member at Oracle Gen2 Cloud and lead teams Block Storage. Prior to that, he started and lead multiple and Transportation optimization infrastructure at Amazc his Masters in Computer Science from Georgia Tech spe performance computing and his Bachelors in Engineerin Science from VJTI Mumbai. Some of my prior talks and Kubecon NA 2019: <https://kccnna19.sched.com/event/cloud-native-machine-learning-data-processing-platform-haytham-abuefutuh-lyft> AWS re:invent 2019: https://www.portal.reinvent.awsevents.com/connect/session_ID=98485&class=popup&srftkn=V300-96FA-M7IV-HGK3-A4E0-BZII Flink Forward 2019: <https://www.ververica.com/resources/flink-forward-sans-2019managing-flink-on-kubernetes-flinkk8soperator> Twi interview: <https://twimlai.com/twiml-talk-343-scalable-ai-maintainable-workflows-at-lyft-with-flyte-w-haytham-abu-ketan-umare/>

SLIDES

Flyte: Cloud Native Machine Learning & Data Processing Platform

Ketan Umare & Matthew Smith / Lyft

#ososummit





VIDEO

SLIDES

Q&A WITH SPEAKER [SPEAKER BIO](#)

Ketan Umare
Software engineer
Lyft

Ketan Umare is a senior staff software engineer at Lyft a the Flyte project. Before Flyte he worked on ETA, routing infrastructure at Lyft. He is also the founder of Flink Kub-operator and contributor to Spark on K8s. Prior to Lyft he founding member at Oracle Gen2 Cloud and lead teams Block Storage. Prior to that, he started and lead multiple and Transportation optimization infrastructure at Amazc his Masters in Computer Science from Georgia Tech spe performance computing and his Bachelors in Engineerin Science from VJTI Mumbai. Some of my prior talks and Kubecon NA 2019: <https://kccnna19.sched.com/event/cloud-native-machine-learning-data-processing-platform-haytham-abuefutuh-lyft> AWS re:invent 2019: https://www.portal.reinvent.awsevents.com/connect/session_ID=98485&class=popup&srftkn=V300-96FA-M7IV-HGK3-A4E0-BZI Flink Forward 2019: <https://www.ververica.com/resources/flink-forward-sans-2019managing-flink-on-kubernetes-flink&operator-Twi> interview: <https://twimlai.com/twiml-talk-343-scalable-ai-maintainable-workflows-at-lyft-with-flyte-w-haytham-abu-ketan-umare/>

Flyte

Production Grade Orchestration for Data and ML

Ketan Umare
Software Engineer
Lyft
[@KetanUmare](#)
[@KetanUmare](#)

Matthew Smith
Software Engineer
Lyft
[@matthewphsmith](#)
[@matthewphsmith](#)





VIDEO



SLIDES

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
Architecture of Flyte
Challenges and learnings while using Kubernetes at scale for stateful applications
A deeper look at developer ergonomics and reasons that shaped some decisions
An overview of the community and ecosystem since open sourcing Flyte
The talk will conclude with a demo.

Agenda



Motivation & Goal

What problem are we trying to solve?

Concepts & Features

Concepts, features and user interface

Architecture & Challenges

Architecture of Flyte and summary of challenges we faced

Case Study

How to use Flyte in the real world

Demo

Everyone loves demos!

Conclusion

Learn more, get involved, & get started





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

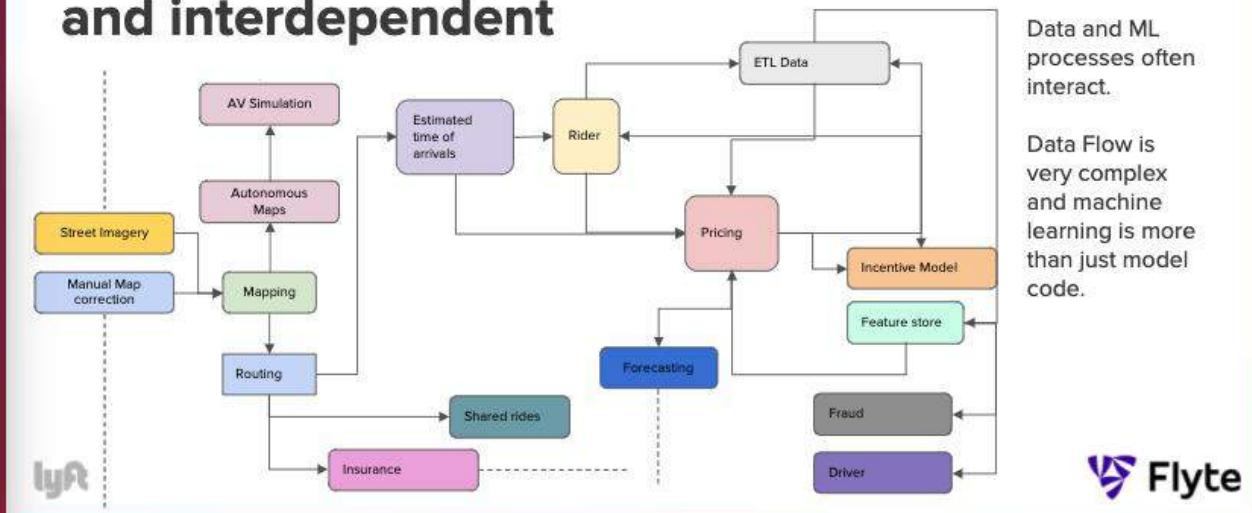
An overview of the community and ecosystem since open-sourcing Flume

open sourcing Flyte
The talk will conclude with a demo

SLIDES

Motivation & Goal

ML & Data services are increasingly complex and interdependent



Data and ML
processes often
interact.

Data Flow is very complex and machine learning is more than just model code.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Motivation & Goal**The parts of the stack we wanted to tackle**

Source: Souley et al. Hidden Technical Debt in Machine Learning Systems





VIDEO



SLIDES

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
Architecture of Flyte
Challenges and learnings while using Kubernetes at scale for stateful applications
A deeper look at developer ergonomics and reasons that shaped some decisions
An overview of the community and ecosystem since open sourcing Flyte
The talk will conclude with a demo.



**Hosted, scalable and serverless
Orchestration Platform**

**Fabric that connects disparate compute
technologies**

Extensible, Observable & shareable

**Integrates best of the breed open source
solutions**

Auditable, Repeatable & Secure





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

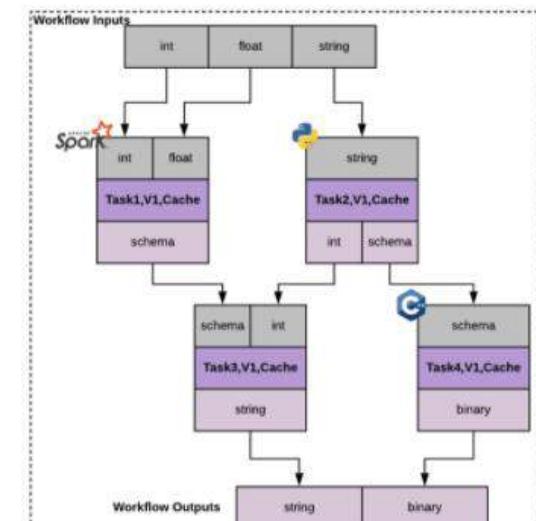
The talk will conclude with a demo.

SLIDES

Concepts & Features

Tasks & Workflows

- Declarative (protobuf)
- Versioned
- Strongly typed interfaces
- Models the flow of Data
- Tasks
 - Arbitrarily complex
 - Encapsulate user code
- Workflows
 - Composable
 - Dynamic
 - DSL





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Concepts & Features**Tasks**

Atomic unit of work & entrypoint to user code. Language agnostic.

Can be executed independently.

```
case class SumTaskInput(a: Long, b: Long)
case class SumTaskOutput(c: Long)
class SumTask extends ... {
    override def
run(input:SumTaskInput):SumTaskOutput = {
    SumTaskOutput(input.a + input.b)
}
https://github.com/spotify/flytekit-java
```



flytekit Scala

Spark Code

```
@inputs(rides=Types.Schema[...],
k=Types.Integer)
@outputs(dest=[Types.String])
@spark_task(spark_conf={"executors":2})
def find_topk_destinations(ctx, spark_ctx,
rides, k, dest):
    ...
    Find the top k destinations for the given set
    of rides ordered by frequency
```

Arbitrary containers

```
edges = SdkRawContainerTask(
    input_data_dir="/inputs",
    output_data_dir="/outputs",
    inputs={"image": Types.Blob, "script": Types.Blob},
    outputs={"edges": Types.Blob},
    image="jjanzic/docker-python3-opencv",
    command=["python", "{{.inputs.script}}",
"/inputs/image", "/outputs/edges"],
)
https://github.com/lyft/flytekit
```





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
 Architecture of Flyte
 Challenges and learnings while using Kubernetes at scale for stateful applications
 A deeper look at developer ergonomics and reasons that shaped some decisions
 An overview of the community and ecosystem since open sourcing Flyte
 The talk will conclude with a demo.

SLIDES

Concepts & Features**Workflows**

Specify the data dependency between tasks (as DAGs). **Composable & declarative!**

Multiple schedules for the same Workflow

Compose with other WF's

```
@workflow_class
class StaticSubWorkflowCaller(object):
    in_ = Input(Types.Integer, default=5,
               help="Input for inner workflow")
    identity_wf_execution =
    IdentityWorkflow(a=outer_a)
    out =
    Output(identity_wf_execution.outputs.task_
          output, sdk_type=Types.Integer)
```

ML Model Train example

```
@workflow_class
class TrainModel(object):
    # Accept inputs
    data = Input(Types.Schema[...])
    hyperparam = Input(Types.Float)
    # Split the dataset
    split = split8020(data=data)

    model = fit_xgboost(
        data=split.train,
        hyperparam=hyperparam)

    pred = eval_xgboost(data=split.val,
                         m=model.outputs.v)

    metrics = compute_metrics(
        data=split.val,
        pred=pred.y_pred)
    # Create outputs
    model = Output(model.outputs.v)
    accuracy = Output(metrics.outputs.acc)
```





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Concepts & Features

Dynamism

Flyte Workflows are **statically defined** and **immutable**.

But, they can contain nodes - that can change the shape dynamically.

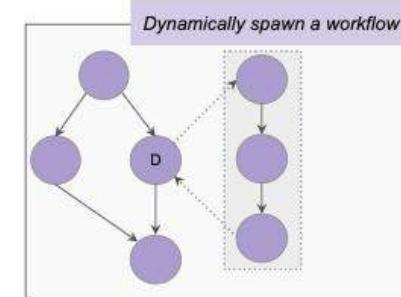
Data parallel jobs, dynamic generation of workflows

(generate logic using the available data) etc

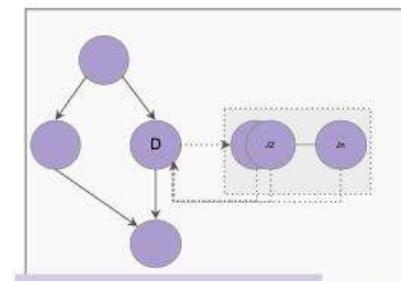
```
@inputs(num=Types.Integer)
@outputs(out=Types.Integer)
@dynamic_task
def sub_wf_yield_task(wf_params, num, out):
    wf_params.logging.info("Running inner task... yielding a sub-workflow")
    identity_wf_execution = IdentityWorkflow(a=num)
    yield identity_wf_execution
    out.set(identity_wf_execution.outputs.task_output)
```



Dynamically spawn a workflow



Dynamically spawn an array of map-only





VIDEO



SLIDES

Concepts & Features

Grouping & Sharing

Projects, Domains & Versions

- Projects offer **logical grouping** of Workflows & Tasks and can be split across one or more repositories, one or more containers
- Domains and Versions provide **CI/CD like semantics** to Workflows & Tasks
 - Users can **push new versions** to production, **rollback** to previous version etc.
 - Users can have workflows in **integration/staging** env
- Domains are **configured globally** for the system (by administrators)

Sharing & Accounting

- Workflows can **refer to tasks and workflows** from other projects
- Executions **accounted/billed** under the **requesters project & domain (Infraspend)**





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Concepts & Features**Shareability: Flytekit Example**

```
Project: ProjectA
@workflow_class
class PipelineA(object):
    in1 = Input(Types.Integer)
    in2 = Input(Types.Integer)
    ...
    out1 = Output(print2.outputs.out)
```

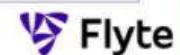
```
Project: ProjectA
@inputs(x=Types.Integer, y=Types.Integer)
@outputs(z=Types.Integer)
@task
def my_model(x, y):
    ...
```

```
Project: ProjectB
@workflow_class
class CompositePipeline(object):
```

```
composed_wf = lps.fetch(
    "ProjectA",
    "Production",
    "PipelineA",
    "1.0.2"
)(in1, in2)
```

```
t1 = local_task(composed_wf.outputs.out)

t2 = tasks.fetch(
    "ProjectA",
    "Production",
    "my_model",
    "2.0.0"
)(x=t1.outputs.x, y=10)
```





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Concepts and Features**DataCatalog: Lineage & Memoization**

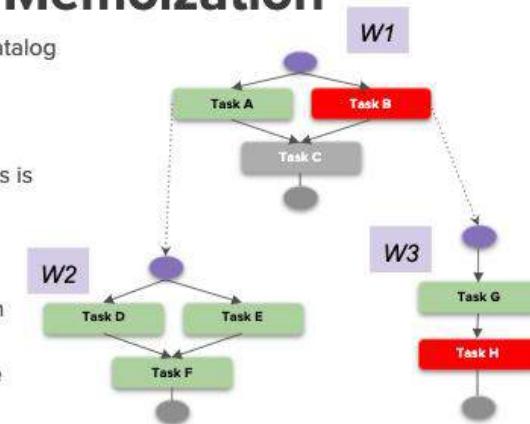
Every task execution in Flyte is **recorded** by default in Catalog Service. This enables Flyte executions to have,

Artifact Lineage

- **Causal dependencies** between data and processes is tracked

Memoization

- Each task execution has a **unique signature**, which includes the input values & version of code
- **Repeated** executions with matching signatures are cached





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Concepts & Features

Serverless

User should only worry about business logic

- specify **resource requirements** - CPU, GPU, Memory, #spark executors etc
- develop **multiple versions of code concurrently**
- **Multi-tenancy** unaware
- Simple **gRPC/REST API** to access all the power
- Language agnostic - **flytekit** (python) and **flytekit-java**, raw containers
- Get **notifications and alerts** for specific events (failures, successes etc)
- **Retrieve results** of past executions

```
@python(gpu_hint=1, cpu=4,  
retries=3, timeout=30s)  
def myFunction():  
    ...
```





VIDEO



ABSTRACT

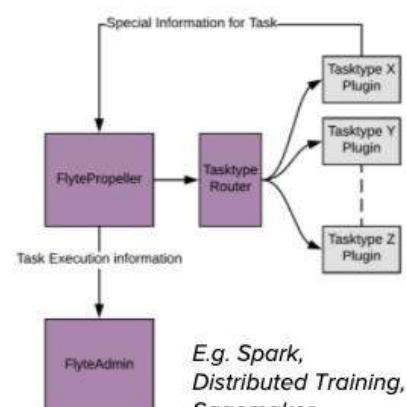
Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
 Architecture of Flyte
 Challenges and learnings while using Kubernetes at scale for stateful applications
 A deeper look at developer ergonomics and reasons that shaped some decisions
 An overview of the community and ecosystem since open sourcing Flyte
 The talk will conclude with a demo.

SLIDES

Concepts & Features

Extensible



```

@sensor_task
def my_test_task(ctx):
    ...
    # e.g. sensor that waits for a hive partition
    # to land. This is added as a contrib.
    ...
    return MyHivePartitionSensor()
  
```

- Container executions are purely from Flyte's POV. (you can write in python, Java etc)
- But this is limited to the implementation SDK
- Backend extensions allow extending capabilities of Flyte.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Architecture & Challenges**Real Production Scale****8.5k Unique Workflows defined****54k Unique Task definitions****1million+ Workflow executions per month****10 million+ task executions per month****40 million+ containers executed per month**

In early January this year, we had a sudden peak and ... (we will get to the story in a bit)



#Workflows per week





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

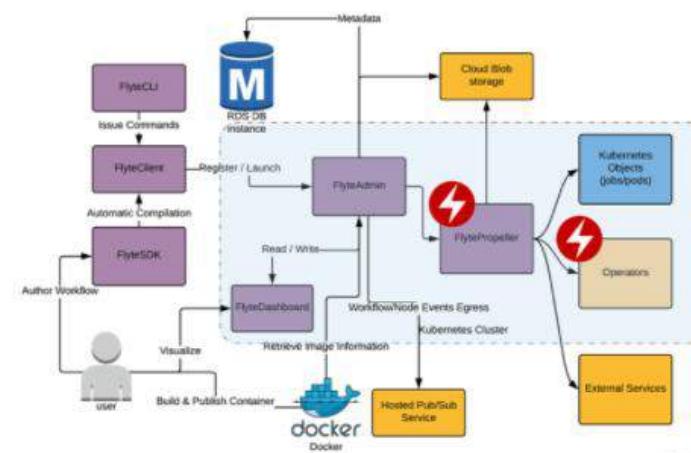
An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Architecture & Challenges**Architecture Overview**

Default: Single Kubernetes cluster with scale-out options to cloud services like AWS Batch.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

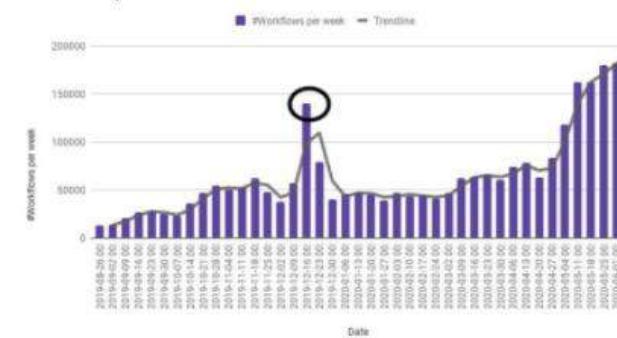
The talk will conclude with a demo.

SLIDES

Architecture & Challenges**Challenges**

- Super-exponential growth, spiked 6x in 2 days
- Flyte had problems
 - Users **lacked visibility**
 - System admins were overwhelmed with operations
 - The cost expenditure ballooned
 - Various scaling problems
 - K8s control plane
 - etcD and K8s Controller
 - K8s Scheduler
 - FlytePropeller is complex (data handling)
- We started diving deeper and optimizing

#Workflows per week





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

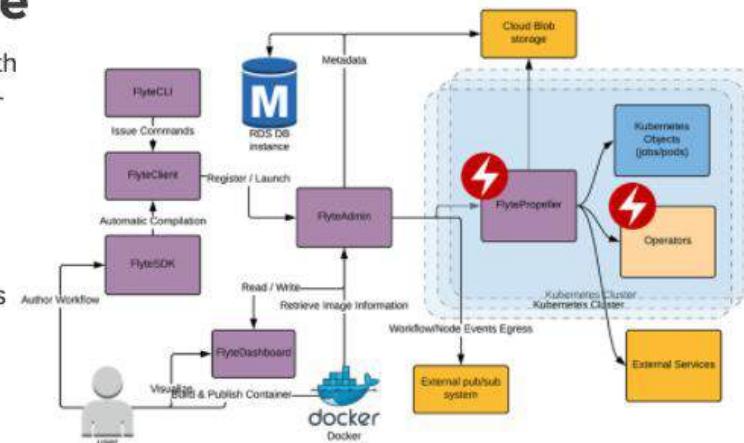
The talk will conclude with a demo.

SLIDES

Architecture & Challenges

Challenge: Scale

- We observed problems with Single Kubernetes cluster -
 - API latency
 - Pod startup issues
 - etcD object size limits
- Single instance of **FlytePropeller*** can run more than **2000** workflows concurrently
 - Informer **caches**
 - Smart control loop short circuits
 - Status **compression**
 - many more!





VIDEO



SLIDES

Architecture & Challenges

Challenge: Multi-tenancy

- Flyte projects are multi-tenancy primitives. Some tenants have larger use-cases as compared to others
- Flyte leverages ResourceQuotas from Kubernetes
 - Flyte Cluster resource controller allows dynamic modification of limits
 - Observability tools show current utilization
 - FlytePropeller backs off intelligently to relieve KubeAPI pressure
- Flyte provides a custom Resource manager on top of Kubernetes Quotas that ensure Global limits and per tenant limits
 - These limits help maintain downstream service
 - Queues to maintain fairness (OSS in progress)
- K8s CRD FairQ (inprogress)





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Architecture & Challenges

Challenge: Visibility

- Automatically generated user dashboard - **grafana template** (oss soon)
- Errors from execution pulled into the UI
- Logs from distributed tasks like **Spark** are pulled into the UI
- Users can **triage** amount of **CPU/memory utilized** by single execution
- **User vs System errors** are clearly separated





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Architecture & Challenges

Challenge: Operations

- Standardized Grafana template for Admins (OSS soon)
- Improved documentation and examples (Work in progress)
- Staged rollouts
- **FlyteAdmin** provides a **flexible routing** system to multiple K8s clusters
 - Allows **isolating** important usecases in different clusters
 - Deployments **bake in lower priority** clusters before proceeding





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Architecture & Challenges

Challenge: Efficiency

- Centralized platform
 - Amortize TCO
 - **Efficiency** multiplier
 - Centralized tooling to visualize costs
 - Utilize **Spot** instances (AWS only)
 - Optimized Cluster Autoscaler (reduced spend by 25%)
 - *Coming soon:* K8s scheduler optimizations



 Flyte



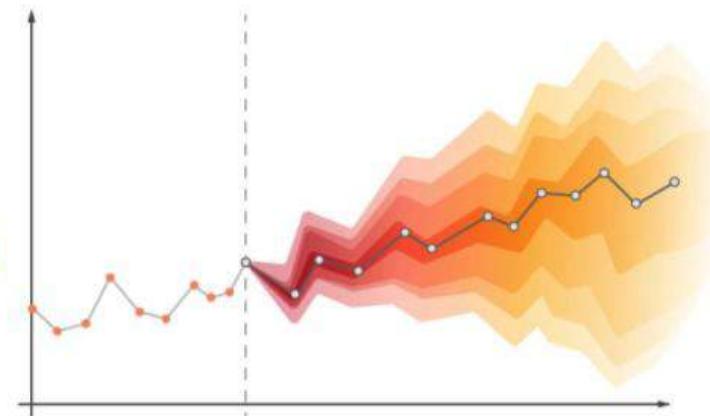
VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
Architecture of Flyte
Challenges and learnings while using Kubernetes at scale for stateful applications
A deeper look at developer ergonomics and reasons that shaped some decisions
An overview of the community and ecosystem since open sourcing Flyte
The talk will conclude with a demo.

SLIDES

**Case Study
Real-time
Forecasting**



VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study

Background: Real-Time Forecasts

- **Business-focused** streaming infrastructure
- **Offline training and tuning**
- **Model health**
 - Backtests
 - In-prod evaluation
- **Workflow-driven deployments**
 - Integration tests
 - Schema evolution
 - Offline data processes
 - Late-arriving trained artifacts





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

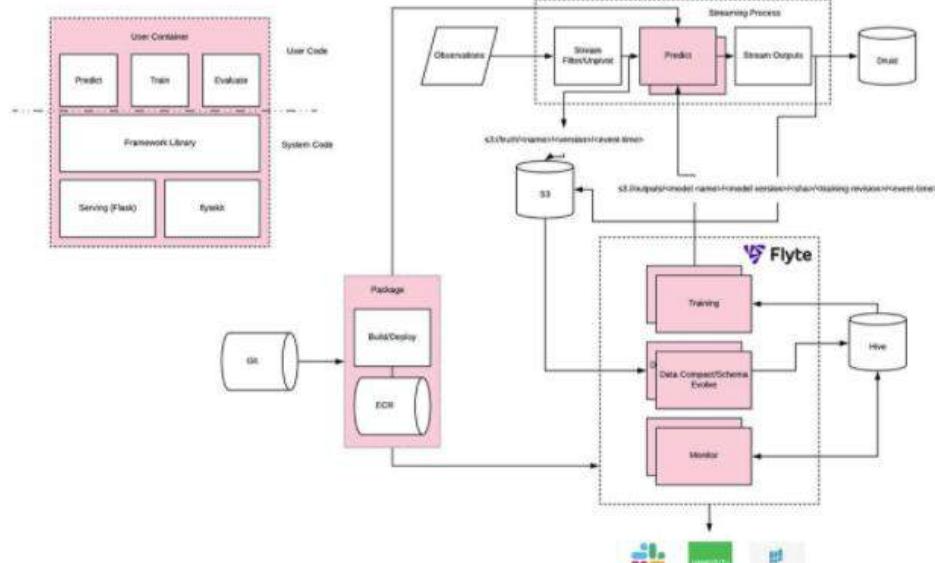
Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study

User Journey

Ideate & Iterate

1. Write business logic
2. Test task locally and remote
3. Orchestrate multiple tasks into a Workflow
4. Execute the workflow
5. Repeat

Productionize

1. Promote a pipeline to production (CI/CD)
2. Create one or more schedules
3. Execute adhoc
4. Monitor and get notified

Retrieve & Replay

1. Retrieve results from executions
2. Identify production errors
3. Replay, reproduce historical artifacts
4. Retrieve artifact lineage





VIDEO



SLIDES

Case Study

Ideate: Write Business Logic

Ideate & Iterate

Productionize

Retrieve & Replay

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

Model Author

```
In [ ]: class TrivialPredictRiders(StreamingModel):
    INPUT_SCHEMA = {
        'ts': datetime.datetime,
        'geohash6': str,
        'rider_cnt': int,
    }

    OUTPUT_SCHEMA = {
        'ts': datetime.datetime,
        'geohash6': str,
        'rider_cnt': float,
    }

    def predict(self, dataframe: pandas.DataFrame) -> pandas.DataFrame:
        # Simply choose the same value as from a day ago in the same place
        dataframe['ts'] += datetime.timedelta(days=1)
        return dataframe

    def train(self, dataframe: pandas.DataFrame) -> dict[str, File]:
        pass
```



Framework Author

```
class StreamingModel(metaclass=abc.ABCMeta):
    @abc.abstractmethod
    def predict(self, dataframe: pandas.DataFrame) -> pandas.DataFrame:
        pass

    def predict_from_flyte(
        self,
        flyte_input: FlyteTypes.Schema
    ) -> FlyteTypes.Schema:
        df = flyte_input.read_to_dataframe()
        self._assert_inputs(df)
        predict_df = self.predict(df)
        self._assert_outputs(predict_df)
        return FlyteTypes.Schema.from_dataframe(predict_df)
```





VIDEO



SLIDES

Case Study

Ideate: Framework Abstracts Flyte

Ideate & Iterate

Productionize

Retrieve & Replay

```
#inputs(model_fqdn=FlyteTypes.String, training_revision=FlyteTypes.String, input_data=FlyteTypes.Schema())
#outputs(output_data=FlyteTypes.Schema())
@python_task
def predict_generic(wf_params, model_fqdn, training_revision, input_data, output_data):
    module_name, attr_name = model_fqdn.rsplit('.', -1)
    model_class = getattr(importlib.import_module(module_name), attr_name)
    model = model_class.materialize(training_revision)
    output_data.set(model.predict_from_flyte(input_data))
```

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study

Ideate: Framework Smooths Tests

Ideate & Iterate

Productionize

Retrieve & Replay

```
from forecastlib.flyte_tasks import predict_generic
from flytekit.common.workflow_execution import SdkWorkflowExecution

class BaseStreamingModel:
    # Continued from before
    def locally_replay_inference(
        self,
        flyte_inference_id: str,
    ) -> pandas.DataFrame:
        # Grab inputs from a failed execution -- links to real production data!
        failed_workflow = SdkWorkflowExecution.fetch(flyte_inference_id)
        inputs = failed_workflow.predict_node.inputs

        # Re-run flyte task in a sandboxed context and return results for analysis
        flyte_output_dict = predict_generic.local_execute(**inputs)
        return flyte_output_dict['output_data'].to_dataframe()
```





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study**Ideate: Rapidly Test Business Logic**

Ideate & Iterate

Productionize

Retrieve & Replay

NODE	STATUS	TYPE	START TIME	DURATION	LOGS
video-model	FAILED	Spark Task	10/7/2019 11:00:18 PM UTC 10/7/2019 11:00:18 PM PDT	1m 53s	View Logs

```
In [ ]: df = my_model.locally_replay_inference(  
        'id_copied_from_ui'  
)
```





VIDEO



SLIDES

Case Study

Ideate: Test/Evaluate at Scale

Ideate & Iterate

Productionize

Retrieve & Replay

Model Author

- Opens Pull Request
- Waits for container build
- User calls test CLI

Framework Author

- Implements generic tasks + workflows
- Tool to register Flyte entities post-build
- Tool to kickoff Flyte test jobs

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.





VIDEO



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study

Ideate: Test/Evaluate at Scale - Framework

Ideate & Iterate

Productionize

Retrieve & Replay

```
#Workflow class
class TestingWorkflow:
    model_fqdn = Input(Types.String)
    start_time = Input(Types.Datetime)
    end_time = Input(Types.Datetime)

    data_fetch_step = DataFetcherWorkflow(
        model_fqdn=model_fqdn,
        start_time=start_time,
        end_time=end_time,
    )
    training_step = TrainingWorkflow(
        model_fqdn=model_fqdn,
        data=data_fetch_step.outputs.training_data,
    )
    predict_step = predict_generic(
        model_fqdn=model_fqdn,
        training_revision=training_step.outputs.training_revision,
        data=data_fetch_step.outputs.test_data,
    )
    results_step = HealthWorkflow(
        model_fqdn=model_fqdn,
        prediction=predict_step.outputs.predictions,
        truth=data_fetch_step.outputs.test_truth,
    )
    scores = Output(results_step.outputs.scores)
```

```
@click.argument('container_sha')
@click.command()
def register_all(container_sha):
    predict_generic.register(container_sha)
    DataFetcherWorkflow.register(container_sha)
    TrainingWorkflow.register(container_sha)
    HealthWorkflow.register(container_sha)
    TestingWorkflow.register(container_sha)
```





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study**Ideate: Test/Evaluate at Scale - Go!**

Ideate & Iterate

Productionize

Retrieve & Replay

Model Author

```
forecast-cli test models.SimpleModel a23d4c --start-date=2019-01-01 --end-date=2020-01-01
```

```
Executing test...
Results:
{'avg_daily_error': 0.03235,
'daily_error': {'FRI': 0.0321,
'MON': 0.0343,
'SAT': 0.0231,
'SUN': 0.0192,
'THR': 0.0282,
'TUE': 0.0423,
'WED': 0.0292},
'daily_error_variance': 0.04248,
'days_analyzed': 365}
```

**Framework Author**

```
from flytekit.common.sdk_launch_plan import SdkLaunchPlan
@click.command()
@click.argument('model_fqdn')
@click.argument('container_sha')
@click.option('start_time', type=click.DateTime(formats=['%Y-%m-%d']))
@click.option('end_time', type=click.DateTime(formats=['%Y-%m-%d']))
def test(model_fqdn, container_sha, start_time, end_time):
    wf = SdkLaunchPlan.fetch('TrainingWorkflow', container_sha)
    execution = wf.execute(
        {
            'model_fqdn': model_fqdn,
            'start_time': start_time,
            'end_time': end_time,
        }
    )
    execution.wait_for_completion()
    click.echo(execution.outputs.scores)
```





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study

Productionize

Ideate & Iterate

Productionize

Retrieve & Replay

Model Author

- Checks-in code
- Waits for integration tests to pass
- Rolls out stage-by-stage
- Monitors dashboards/alarms

**Framework Author**

- Integration test harness
- Release tools
 - Automated training
 - Offline-data compaction
 - Health monitoring
- Programmatic alarms and dashboards





VIDEO



SLIDES

Case Study

Productionize: Integration

Ideate & Iterate

Productionize

Retrieve & Replay

```
def integration_test(model_fqdn: str, model_config: dict[str, Any], container_sha: str) -> None:
    start_time = datetime.utcnow() - model_config['integration_window_lag']
    end_time = start_time + model_config['integration_window_length']

    wf = TestingWorkflow.fetch(container_sha)
    execution = wf.execute(
        {
            'model_fqdn': model_fqdn,
            'start_time': start_time,
            'end_time': end_time,
        }
    )
    execution.wait_for_completion()

    assert _scores_pass_model_config(execution.outputs.scores, model_config)
```

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.





VIDEO



SLIDES

Case Study

Productionize: Release

Ideate & Iterate

Productionize

Retrieve & Replay

```
from flytekit.common.sdk_workflow import SdkWorkflow
def release(model_fqdn: str, model_config: dict[str, Any], container_sha: str) -> None:
    training_workflow = SdkWorkflow.fetch('TrainingWorkflow', container_sha)
    for training_config in model_config['training_pipelines']:
        lp = training_workflow.create_launch_plan(
            schedule=schedules.CronSchedule(training_config['cron_expr']),
            fixed_inputs=training_config['params'],
        )
        lp.register(training_config['name'], container_sha)
        lp.update_launch_plan(LaunchPlanState.ACTIVE)

    offline_data_organizer = SdkWorkflow.fetch('OfflineDataOrganizer', container_sha)
    lp = offline_data_organizer.create_launch_plan(
        schedule=schedules.CronSchedule(model_config['offline_data_latency_cron']),
        fixed_inputs=model_config['data_schema_params'],
    )
    lp.update_launch_plan(LaunchPlanState.ACTIVE)

    # ... additional offline processes for health checking, cron-based signals, etc.
```



ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
Architecture of Flyte
Challenges and learnings while using Kubernetes at scale for stateful applications
A deeper look at developer ergonomics and reasons that shaped some decisions
An overview of the community and ecosystem since open sourcing Flyte
The talk will conclude with a demo.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

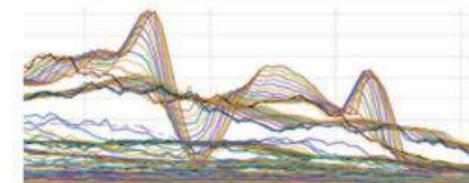
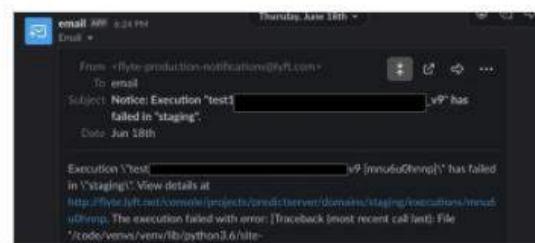
SLIDES

Case Study**Productionize: Monitor**

Ideate & Iterate

Productionize

Retrieve & Replay





VIDEO



SLIDES

Case Study

Retrieve and Replay

Ideate & Iterate

Productionize

Retrieve & Replay

Model Author

- Notices anomaly in behavior
- Fetches exact input data
- If desired, fetches exact container
- Go back to 'Ideate & Iterate'

Framework Author

- One-click tools for common debug paths.
 - Retry inference
 - Compare inferences over time

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.





VIDEO

**ABSTRACT**

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

Case Study**Retrieve and Replay: Isolated Failures**

Ideate & Iterate

Productionize

Retrieve & Replay

Model Author

```
forecast-cli debug failure --inference-id c21a8c
```

```
In [1]: # THE CODE IN THIS CELL IS GENERATED!
from models import TrivialPredictRiders

TRAINING_REVISION = 'f32elc'
INFERENCE_ID = 'c21a8c'

model = TrivialPredictRiders.create_from_trained_artifacts(
    TRAINING_REVISION
)
input_dataframe = model.fetch_inference_inputs(INFERENCE_ID)
output_dataframe = model.fetch_inference_outputs(INFERENCE_ID)
```

**Framework Author**

```
from flytekit.common.workflow_execution import SdkWorkflowExecution
from flytekit.common.launch_plan import SdkLaunchPlan
from flytekit.common.tasks.task import SdkTask
from flytekit.common.workflow import SdkWorkflow
```

```
def get_debug_params(inference_id: str) -> dict[str, Any]:
    flyte_execution_id = _map_inference_to_flyte_id(inference_id)
    execution = SdkWorkflowExecution.fetch(flyte_execution_id)
    executed_launch_plan = SdkLaunchPlan.fetch(
        execution.spec.launch_plan
    )
    executed_workflow = SdkWorkflow.fetch(
        executed_launch_plan.spec.workflow_id
    )
    executed_task = SdkTask.fetch(
        executed_launch_plan.spec.workflow_id
    )
    executed_workflow.nodes['predict-node'].task_node.reference_id
    )

    model_fpds = execution.inputs.model_fpds
    training_revision = ...
```

```
    return {
        'model_fpds': execution.inputs.model_fpds,
        'docker_image': executed_task.container.image,
        'training_revision': execution.inputs.training_revision,
    }
```





VIDEO



SLIDES

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.



Demo



Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

In [25]: `from image import filter_edges`

In []: `inp = "image.png" # This should be a path to an image that is available locally
out = "edges.png" # This is the path to where we want to create the image
filter_edges(inp, out)
Invoke the helper method that displays the image in Jupyter
display_images([inp, out])`

Step I: Run the filter_edges example within Jupyter Notebook First

Step II: Create a Flyte Task for this function

In this case we will use an SdkRawContainerTask. A raw container task is essentially a container task, where we tell Flyte, that this container does not have flytekit. So all the inputs and outputs should be auto-mounted and uploaded. The task can use an open source container `docker.io/jjanzic/docker-python3-opencv`. This container has python and OpenCV already installed.

Hot load the code

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 O

In []:

```
from flytekit.common.tasks.raw_container import SdkRawContainerTask
from flytekit.sdk.types import Types

edges = SdkRawContainerTask(
    input_data_dir="/inputs",
    output_data_dir="/outputs",
    inputs={"image": Types.Blob, "script": Types.Blob},
    outputs={"edges": Types.Blob},
    image="docker.io/jjanzic/docker-python3-opencv",
    command=["python", "/inputs/script", "/inputs/image", "/outputs/edges"]
)
```

Step IIIa: Just for this excercise

To make the dynamic loading of the script work, we have to **upload the script to some s3 bucket**. Since we are testing this locally, I have originally created an s3 client that points to Flyte installed minio. Upload the code there too

In []:

```
e/image.py"
h)
photographysimplified.com/wp-content/uploads/2019/06/How-to-get-sharp-images
```

refresh back forward search user help

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

command=["python", "/inputs/script", "/inputs/image", "/outputs/edges"]

Step IIIa: Just for this excercise

To make the dynamic loading of the script work, we have to **upload the script to some s3 bucket**. Since we are testing this locally, I have originally created an s3 client that points to Flyte installed minio. Upload the code there too

```
In [ ]: le/image.py
h)
photographysimplified.com/wp-content/uploads/2019/06/How-to-get-sharp-images
```

Step III: Launch an execution for the task

This creates an execution of just the task. Remember in Flyte, Task is a standalone top level entity, so you can execute it

```
In [ ]: exc = edges.register_and_launch("flyteexamples", "development", inputs={"in
```

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

Step III: Launch an execution for the task

This creates an execution of just the task. Remember in Flyte, Task is a standalone top level entity, so you can execute it

```
In [ ]: exc = edges.register_and_launch("flyteexamples", "development", inputs={"ir": ...})
print_console_url(exc)
```

```
In [ ]: exc.wait_for_completion()
```

Step IV: Optional. Create a Workflow

Ofcourse you can use this task in a workflow. We are creating a trivial workflow in this case that has only one task

```
In [ ]: from flytekit.sdk.workflow import workflow_class, Input, Output
@workflow_class
class EdgeDetector(object):
    script = Input(Type.Blob)
    image = Input(Type.Blob)
    edge_task = edges(script=script, image=image)
```



SCREEN SHARE

Open Source Summit + En Home Page - Select or create raw-container - Jupyter Note Flyte Console sagemaker-hpo-demo - Jupyter How-to-get-sharp-images-BI

localhost:30081/console/projects/flyteexamples/domains/development/executions/rI90yea052

← RUNNING flyteexamples/d... rI90yea052 Domain development Version 744ea827a... Time 7/1/2020 8:51:57 PM UTC Duration View Inputs & Outputs Terminate

Nodes Graph

Status Start Time Duration

NOD	STATUS	TYPE	START TIME	DURATION Queued Time	LOGS
mair	RUNNING	Unknown Task	7/1/2020 8:51:58 PM UTC 7/1/2020 1:51:58 PM PDT	2s	View Logs



Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.nxp

SCREEN SHARE

jupyter sagemaker-hpo-demo (autosaved)

Logout

Trusted | Python 3.7.4 64-bit ('flytekit': virtualenv) O

File Edit View Insert Cell Kernel Widgets Help

In []:

```
from flytekit.configuration import set_flyte_config_file, platform
set_flyte_config_file("/Users/kumare/.ssh/notebook-staging.config")
#set_flyte_config_file("notebook.config")

print("Connected to {}".format(platform.URL.get()))

def print_console_url(exc):
    print("http://{}{}".format(console.projects[{}].domains[{}].executions[{}].format(p[0], p[1], p[2], p[3]))
```

Step I: Algorithm for training

XGBoost

In []:

```
from flytekit.sdk.tasks import inputs
from flytekit.sdk.types import Types
from flytekit.sdk.workflow import workflow_class, Input, Output
from flytekit.common.tasks.sagemaker import training_job_task, hpo_job_task
from flytekit.models.sagemaker import training_job as training_job_models,
from flytekit.sdk.sagemaker import types as _sdk_sagemaker_types
```

Screen sharing controls: Refresh, Stop, Share, Chat, User, Document, Video.

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter sagemaker-hpo-demo (unsaved changes)

Logout

Trusted Python 3.7.4 64-bit ('flytekit': virtualenv)

File Edit View Insert Cell Kernel Widgets Help

In [12]:

```
from flytekit.sdk.tasks import inputs
from flytekit.sdk.types import Types
from flytekit.sdk.workflow import workflow_class, Input, Output
from flytekit.common.tasks.sagemaker import training_job_task, hpo_job_task
from flytekit.models.sagemaker import training_job as training_job_models,
from flytekit.sdk.sagemaker import types as _sdk_sagemaker_types

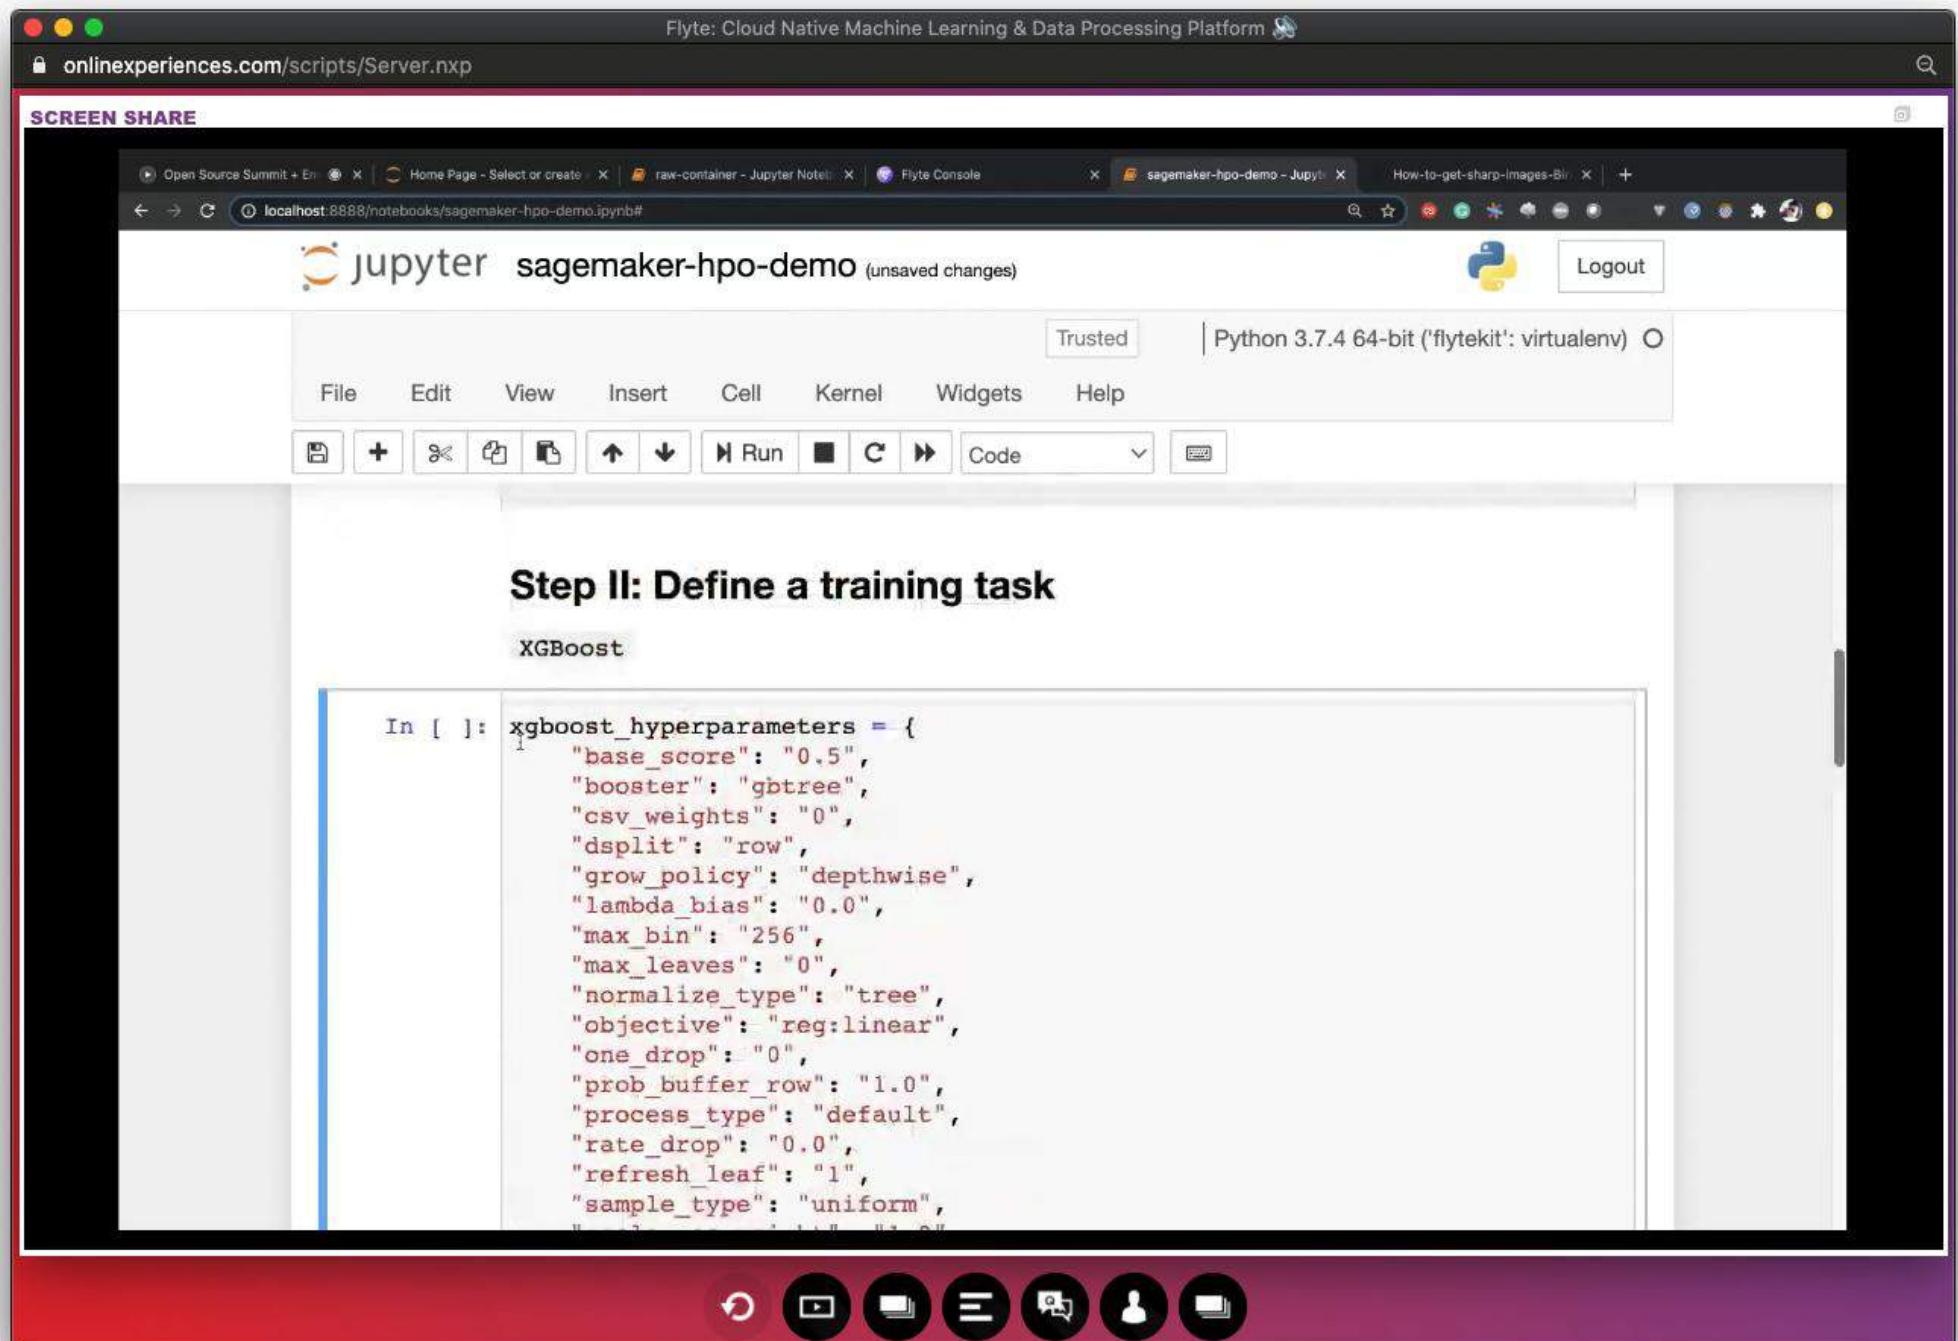
alg_spec = training_job_models.AlgorithmSpecification(
    input_mode=_sdk_sagemaker_types.InputMode.FILE,
    algorithm_name=_sdk_sagemaker_types.AlgorithmName.XGBOOST,
    algorithm_version="0.72",
    metric_definitions=[training_job_models.MetricDefinition(name="Minimize")
)
```

Step II: Define a training task

XGBoost

In []:

```
xgboost_hyperparameters = {
    "base_score": "0.5",
    "n_estimators": "100000"
}
```



Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter sagemaker-hpo-demo (unsaved changes)

Logout

Trusted | Python 3.7.4 64-bit ('flytekit': virtualenv) O

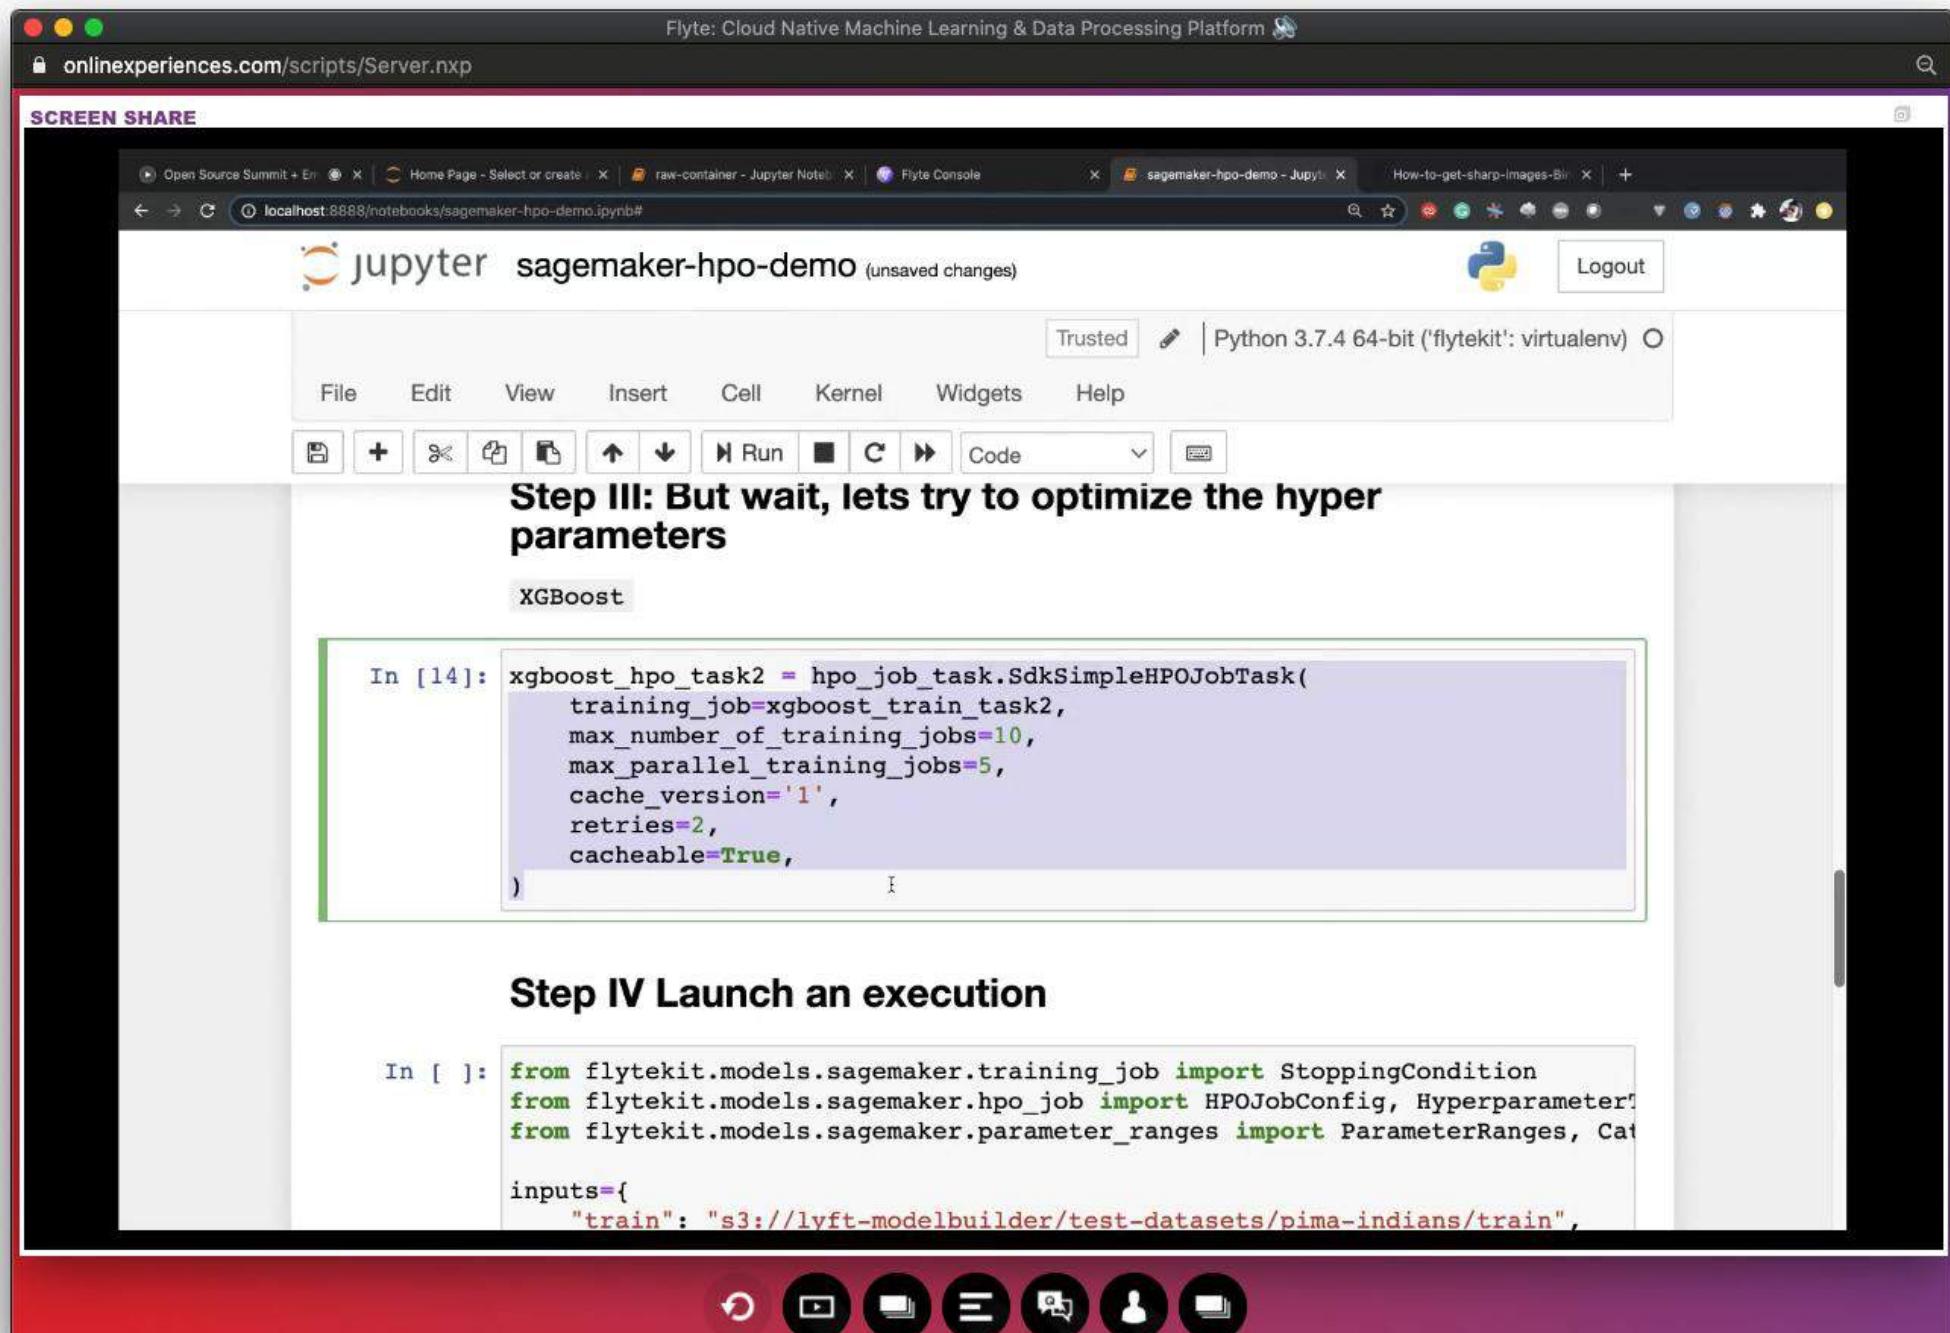
File Edit View Insert Cell Kernel Widgets Help

Code

```
    "refresh_leaf": "1",
    "sample_type": "uniform",
    "scale_pos_weight": "1.0",
    "silent": "0",
    "skip_drop": "0.0",
    "tree_method": "auto",
    "tweedie_variance_power": "1.5",
    "updater": "grow_colmaker,prune",
}

xgboost_train_task2 = training_job_task.SdkSimpleTrainingJobTask(
    training_job_config=training_job_models.TrainingJobConfig(
        instance_type="ml.m4.xlarge",
        instance_count=1,
        volume_size_in_gb=2,
),
    algorithm_specification=alg_spec,
    cache_version='1',
    cacheable=True,
)
```

Step III: But wait, lets try to optimize the hyper parameters





SCREEN SHARE



Open Source Summit + Home Page - Select or creat raw-container - Jupyter Note Flyte Console sagemaker-hpo-demo - Jup Flyte Console How-to-get-sharp-images... +
← → C flyte-staging.lyft.net/console/projects/flyteexamples/domains/development/executions/elzaoqgt2t

UNKNOWN flyteexamples/development/Bytegen_main_xgboost_hpo_task2 elzaoqgt2t Domain development Version fbeafc95a038a12d7b93a4b0fb01beb7 Time Duration View Inputs & Outputs Terminate

Nodes Graph

Status Start Time Duration

NODE	STATUS	TYPE	START TIME	DURATION	LOGS
------	--------	------	------------	----------	------

No executions found.





SCREEN SHARE

Open Source Summit + | Home Page - Select or crea... | raw-container - Jupyter Note... | Flyte Console | sageMaker-hpo-demo - Jupyter Note... | Flyte Console | How-to-get-sharp-images... | +
localhost:30081/console/projects/flyteexamples/domains/development/executions/rI90yea052

SUCCEEDED flyteexamples/d... rI90yea052 Domain development Version 744ea827a... Time 7/1/2020 8:51:57 PM UTC Duration 16s View Inputs & Outputs Relaunch

Nodes Graph

Status Start Time Duration

NOD	STATUS	TYPE	START TIME	DURATION Queued Time	LOGS
> mair	SUCCEEDED	Unknown Task	7/1/2020 8:51:58 PM UTC 7/1/2020 1:51:58 PM PDT	14s	View Logs





SCREEN SHARE



Open Source Summit +	Home Page - Select or create	raw-container - Jupyter Note	Flyte Console	sagemaker-hpo-demo - Jup	Flyte Console	How-to-get-sharp-images	+
localhost:30081/console/projects/flyteexamples/domains/development/executions/rI90yea052							

SUCCEEDED

flyteexamples/d... rI90yea052 Domain development Version 744ea827a... Time 7/1/2020 8:51:57 PM UTC Duration 16s View Inputs & Outputs Relaunch

Nodes Graph

Status Start Time Duration

NOD	STATUS	TYPE	START TIME
mair	SUCCEEDED	Unknown Task	7/1/2020 8:51:58 PM UTC 7/1/2020 1:51:58 PM PDT

mainedges

SUCCEEDED

TYPE
Unknown Task

Executions Inputs Outputs Task

main.edges

succeeded

Logs

Kubernetes Logs (User)

started (unknown)

run time (unknown)





SCREEN SHARE



Open Source Summit + Home Page - Select or creat raw-container - Jupyter Note Flyte Console sagemaker-hpo-demo - Jup Flyte Console How-to-get-sharp-images- Relaunch

localhost:30081/console/projects/flyteexamples/domains/development/executions/r190yea052



SUCCEEDED

flyteexamples/d... Domain Version
rl90yea052 development 744ea827a...

Time Duration
7/1/2020 8:51:57 PM UTC 16s

View Inputs & Outputs

Relaunch

rl90yea052



Inputs

Outputs

image:
type: single
uri: https://www.naturephotographysimplified.com/wp-content/uploads/2019/06/How-to-get-sharp-images-Birds-in-flight-Bharatpur-Bird-Sanctuary-bird-Photography-by-Prathap-DK-bronze-winged-jacana-Greater-Spotted-Eagle-750x500.jpg
script:
type: single
uri: s3://my-s3-bucket/code/image.py

Task





SCREEN SHARE



Open Source Summit +	Home Page - Select or create	raw-container - Jupyter Note	Flyte Console	sagemaker-hpo-demo - Jup	Flyte Console	How-to-get-sharp-images	+	
←	SUCCEEDED	flyteexamples/d... rl90yea052	Domain development	Version 744ea827a...	Time 7/1/2020 8:51:57 PM UTC	Duration 16s	View Inputs & Outputs	Relaunch

rl90yea052

[Inputs](#) [Outputs](#)

```
edges:  
  type: single  
  uri: s3://my-s3-bucket/zh/r190yea052-mainedges-0/edges
```

> main

Task



Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (autosaved)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

Step III: Launch an execution for the task

This creates an execution of just the task. Remember in Flyte, Task is a standalone top level entity, so you can execute it

```
In [29]: exc = edges.register_and_launch("flyteexamples", "development", inputs={"in": "http://localhost:30081/console/projects/flyteexamples/domains/development/executions/r190yea052"}  
print_console_url(exc)
```

<http://localhost:30081/console/projects/flyteexamples/domains/development/executions/r190yea052>

```
In [ ]: exc.wait_for_completion()
```

Step IV: Optional. Create a Workflow

Ofcourse you can use this task in a workflow. We are creating a trivial workflow in this case that has only one task

```
In [ ]: from flytekit.sdk.workflow import workflow_class, Input, Output  
@workflow_class  
class EdgeDetector(object):
```

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.npx

SCREEN SHARE

jupyter raw-container (autosaved)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 O

In []:

```
from flytekit.sdk.workflow import workflow_class, Input, Output
@workflow_class
class EdgeDetector(object):
    script = Input(Types.Blob)
    image = Input(Types.Blob)
    edge_task = edges(script=script, image=image)
    out = Output(edge_task.outputs.edges, sdk_type=Types.Blob)

EdgeDetector_lp = EdgeDetector.create_launch_plan()
```

Step V: Optional. Register and execute the workflow

To make the dynamic loading of the script work, we have to upload the script to some s3 bucket. Since we are testing this locally, I have originally created an s3 client that points to Flyte installed minio. Upload the code there too

`s3.Bucket('my-s3-bucket').upload_file('image.py', 'code/image.py')`

In []:

```
edges.register(name="EdgeDetectorFunc", project="flyteexamples", domain="de
EdgeDetector.register(name="EdgeDetector", project="flyteexamples", domain=
EdgeDetector_lp.register(name="EdgeDetector", project="flyteexamples", doma
```

refresh back forward search user help

Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.nxp

SCREEN SHARE

jupyter raw-container (unsaved changes)

Logout

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3

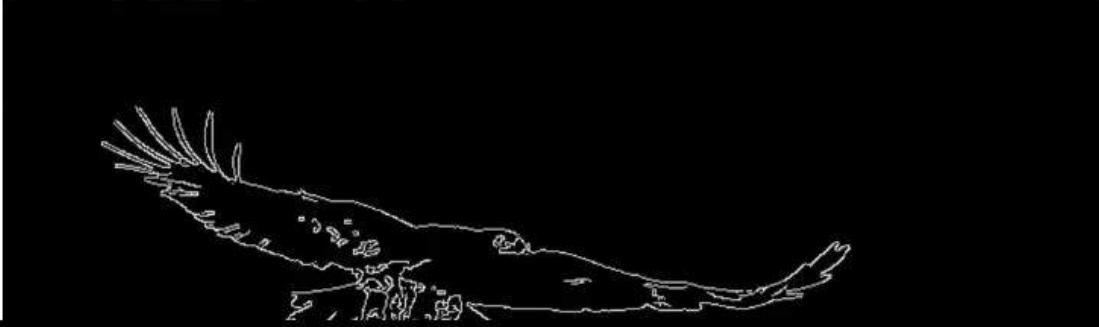
In []: exc = EdgeDetector_lp.execute("flyteexamples", "development", inputs={"image": "https://raw.githubusercontent.com/flyteorg/flyte-examples/main/assets/zebra.png"}, print_console_url(exc)

Step VI: Visualize the results

You can retrieve the results and visualize them here

In [30]: key="s3://my-s3-bucket/zh/r190yea052-mainedges-0/edges" I
download_file(key, "edges.png")

In [31]: display_images(["edges.png"])

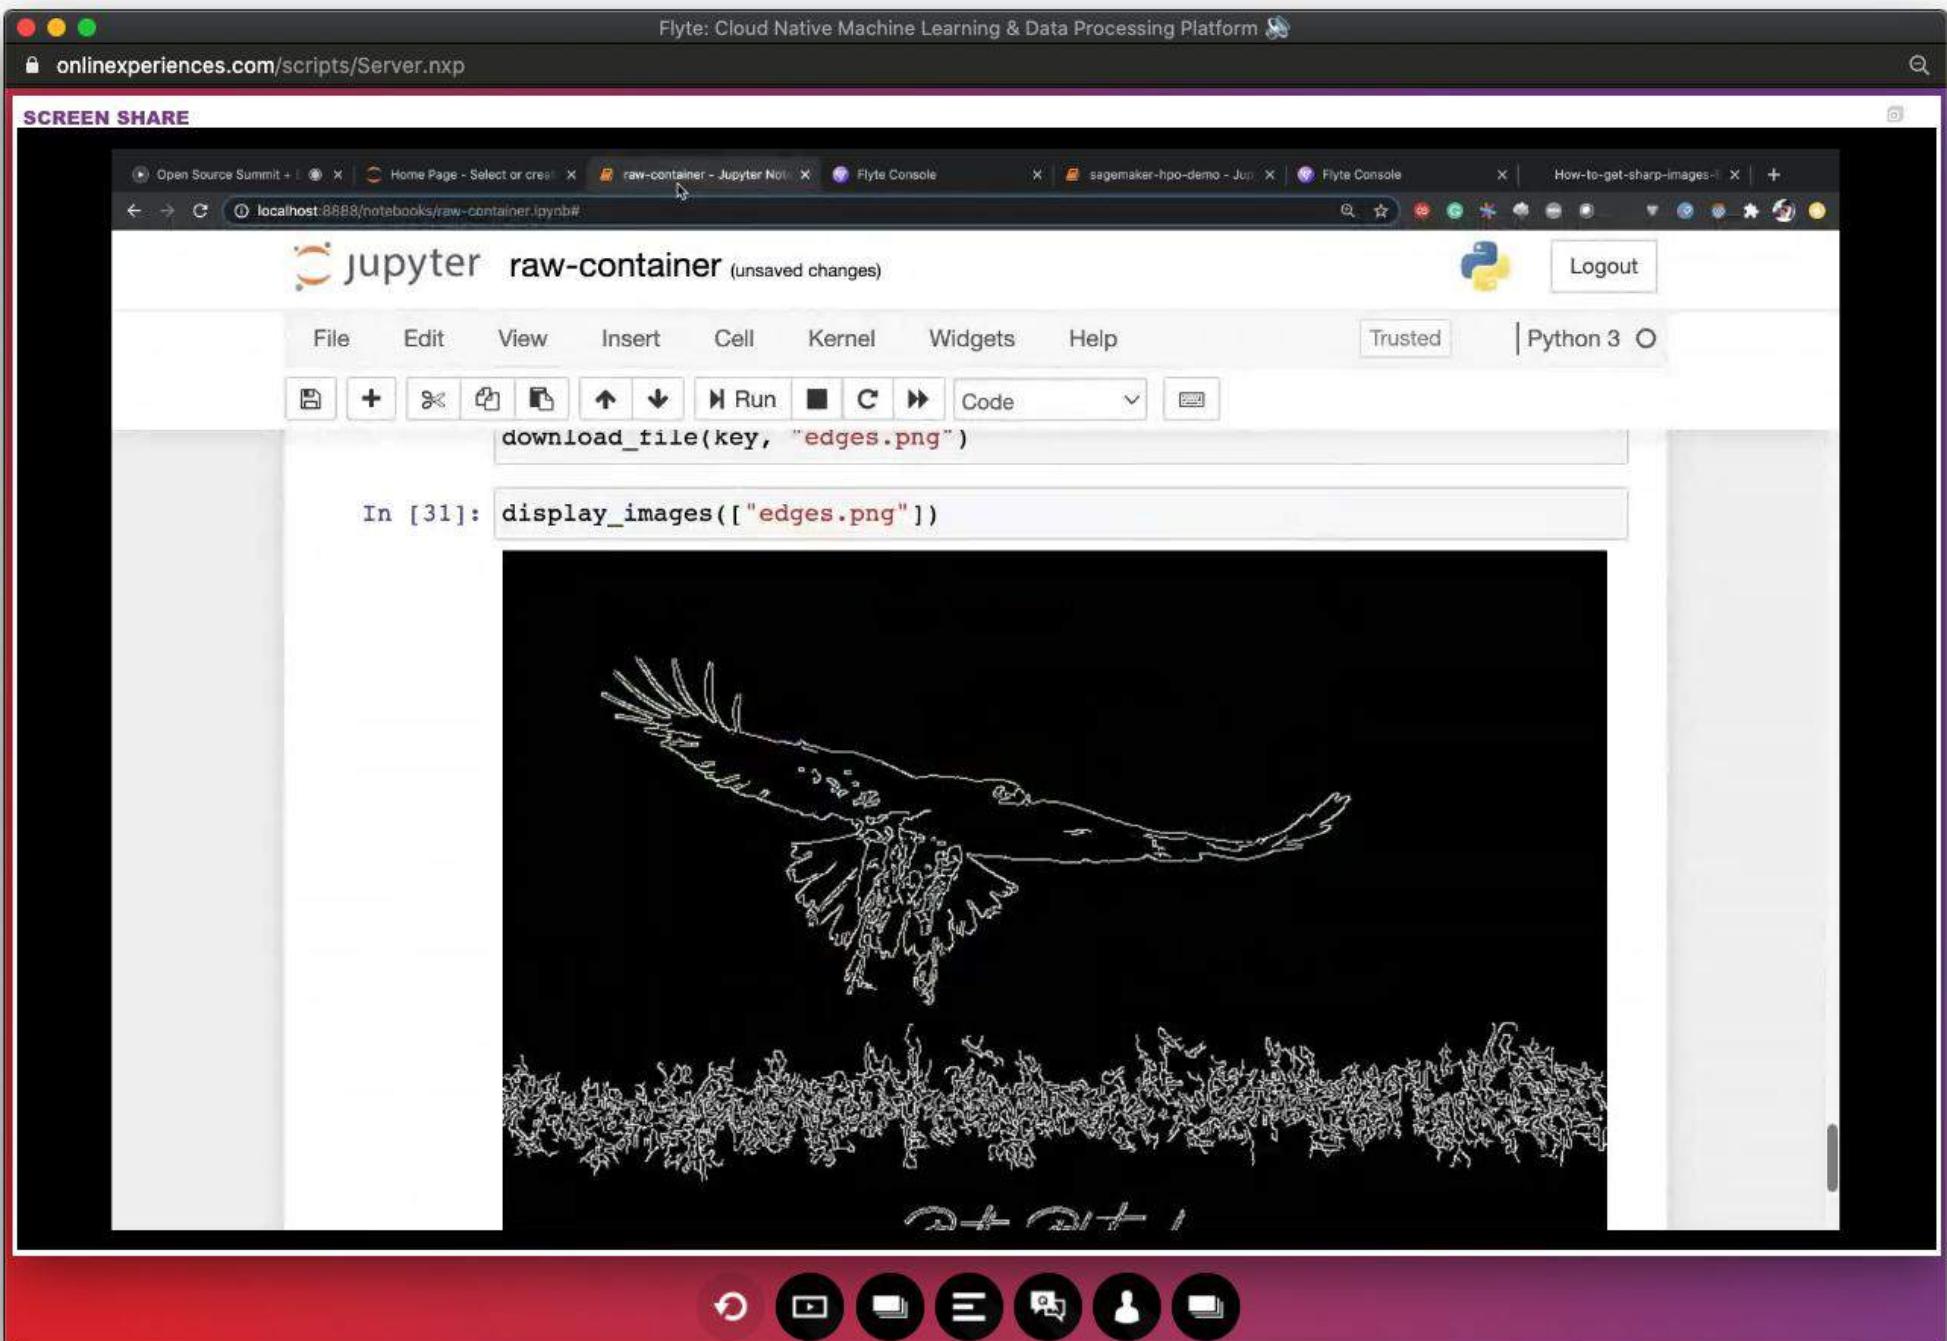


Red navigation icons at the bottom include: back, forward, search, refresh, and user profile.

SCREEN SHARE

Open Source Summit | Home Page - Select or crea... raw-container - Jupyter Note... Flyte Console sagemaker-hpo-demo - Jup... Flyte Console How-to-get-sharp-images... ↗







SCREEN SHARE

NODE	STATUS	TYPE	START TIME	DURATION Queued Time	LOGS
mainxgboosthpotask2 _main_xgboost_hpo_task2	SUCCEEDED	Unknown Task			View Logs



SCREEN SHARE

Open Source Summit + | Home Page - Select or create | raw-container - Jupyter Notebooks | Flyte Console | sageMaker-hpo-demo - Jupyter Notebooks | Flyte Console | How-to-get-sharp-images | +

← SUCCEEDED flyteexamples/development/flytegen_main_xgboost_hpo_task2 Domain development Version fbeafc95a038a12d7b93a4b0fb01b... Time 7/1/2020 8:53:22 PM UTC Duration 42s View Inputs & Outputs Relaunch

Nodes Graph

Status Start Time Duration

NODE	STATUS	TYPE	START TIME
mainxgboosthpotask2 _main_xgboost_hpo_task2	SUCCEEDED	Unknown Task	

mainxgboosthpotask2

_main_xgboost_hpo_task2

SUCCEEDED

TYPE Unknown Task

Executions Inputs Outputs Task

```
hpo_job_config:  
  tag: pb_type=flyteidl.plugins.sagemaker.hpo_  
        job_pb2.HPOJobConfig  
  (binary data not shown)  
  static_hyperparameters: {  
    base_score: 0.5  
    booster: gbt  
    csv_weights: 0  
    dsplit: row  
    grow_policy: depthwise  
    lambda_bias: 0.0  
    max_bin: 256  
    max_leaves: 0  
    normalize_type: tree  
    objective: reg:linear  
    one_drop: 0  
    prob_buffer_row: 1.0  
    process_type: default  
    rate_drop: 0.0  
    refresh_leaf: 1  
    sample_type: uniform  
    scale_pos_weight: 1.0  
    silent: 0  
    tree_limit: 0  
  }
```

刷新 前进 后退 搜索 书签 全屏



SCREEN SHARE

Open Source Summit	Home Page - Select or create	raw-container - Jupyter Notebooks	Flyte Console	sagemaker-hpo-demo - Jupyter Notebooks	Flyte Console	How-to-get-sharp-images	+
flyte-staging.lyft.net/console/projects/flyteexamples/domains/development/executions/elzaoqgt2t	SUCCEEDED	flyteexamples/development/flytegen_main_xgboost_hpo_task2	Domain development	Version fbeafc95a038a12d7b93a4b0fb01b...	Time 7/1/2020 8:53:22 PM UTC	Duration 42s	View Inputs & Outputs Relaunch

Nodes

Graph

Status

Start Time

Duration

NODE

STATUS

TYPE

START TIME

mainxgboosthpotask2
_main_xgboost_hpo_task2

SUCCEEDED

Unknown Task

mainxgboosthpotask2

_main_xgboost_hpo_task2

SUCCEEDED

TYPE

Unknown Task

Executions

Inputs

Outputs

Task

```
model:  
  type: single  
  uri: s3://lyft-modelbuilder/metadata/propeller/staging/flyteexamples-development-ydr3ty28w7/mainxgboosthpotask2/data/0/hpo_outputs/ydr3ty28w7-006-cef6fee4/output/model.tar.gz
```



Flyte: Cloud Native Machine Learning & Data Processing Platform

onlinexperiences.com/scripts/Server.nxp

SCREEN SHARE

jupyter sagemaker-hpo-demo (unsaved changes)

Logout

Trusted Python 3.7.4 64-bit ('flytekit': virtualenv)

File Edit View Insert Cell Kernel Widgets Help

In [17]: `exc.sync()`

In [18]: `m = exc.outputs['model']`

In [19]: `m.uri`

Out[19]: `'s3://lyft-modelbuilder/metadata/propeller/staging/flyteexamples-development-ydr3ty28w7/mainxgboosthpotask2/data/0/hpo_outputs/ydr3ty28w7-006-cef6fee4/output/model.tar.gz'`

Step V

Retrieve results from the execution and use them locally to run predictions?

In []: `m.download("/tmp/model", overwrite=True)`

In []:



VIDEO



SLIDES

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.



Demo





ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts

Architecture of Flyte

Challenges and learnings while using Kubernetes at scale for stateful applications

A deeper look at developer ergonomics and reasons that shaped some decisions

An overview of the community and ecosystem since open sourcing Flyte

The talk will conclude with a demo.

SLIDES

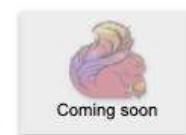
Conclusion

Ecosystem



TensorFlow

Coming soon





VIDEO



SLIDES

Conclusion

What's Next

Flyte is constantly evolving and new features are coming soon like,

- Flytekit python Enhancement & Flytekit JAVA primetime ready
- Richer **data catalog visualization**
- **UI** improvements
- **Reactive workflows** (respond to data publication events)
- Better documentation and more examples
- Faster getting started

To find more details **visit our docs and the Roadmap section**. Also join our fledgeling community and help us shape the future of Flyte. We appreciate contributions and suggestions.





VIDEO



SLIDES

Thanks!
Learn more, get started & keep in
touch at [Flyte.org](https://flyte.org)

ABSTRACT

Flyte is the backbone for large-scale Machine Learning and Data Processing (ETL) pipelines at Lyft. It is used across business critical applications ranging from ETA, Pricing, Mapping, Autonomous etc. At its core it is a Kubernetes native workflow engine that executes 10M+ containers per month as part of thousands of workflows. We introduced Flyte at Kubecon 2019. This talk will build on the base talk and dive deeper into the architecture and the developer experience. Thus the talk will help existing users of Flyte and new users as well.

Quick overview of Flyte concepts
Architecture of Flyte
Challenges and learnings while using Kubernetes at scale for stateful applications
A deeper look at developer ergonomics and reasons that shaped some decisions
An overview of the community and ecosystem since open sourcing Flyte
The talk will conclude with a demo.

@KetanUmare
 @KetanUmare

@matthewphsmith
 @matthewphsmith

